A Project Report

On

AI - POWERED EARLY DETECTION OF CHILD MALNUTRITION

BY

Name	ID Number
Mehak Mahajan	2022A3PS0585H
Khush Bhuta	2022A7PS1333H
Vaishnu Kanna	2022B3A71608H
Utkarsh Singhal	2022A7PS1334H

Under the supervision of

Prof. Subhrakanta Panda

SUBMITTED IN PARTIAL FULFILLMENT OF THE REQUIREMENTS OF

CS F266: STUDY PROJECT

BIRLA INSTITUTE OF TECHNOLOGY AND SCIENCE PILANI (RAJASTHAN)

HYDERABAD CAMPUS

(DECEMBER 2025)

ACKNOWLEDGMENTS

We would like to acknowledge the support and express my sincere gratitude to **Prof. Subhrakanta Panda**, Department of **Computer Science and Information Systems** at **Birla Institute of Technology and Science Pilani, Hyderabad campus** for overseeing this project. We thank him for his constant guidance, encouragement and valuable insights throughout the project.

Birla Institute of Technology and Science-Pilani,

Hyderabad Campus

Certificate

This is to certify that the project report entitled "**AI POWERED EARLY DETECTION OF CHILD MALNUTRITION**" submitted by Ms. Mehak Mahajan (ID No. 2022A3PS0585H), Mr. Khush Bhuta (ID No. 2022A7PS1333H), Ms. Vaishnu Kanna (ID No. 2022B3A71608H), Mr. Utkarsh Singhal (ID No. 2022A7PS1334H) in partial fulfillment of the requirements of the course CS F266, Study Project Course, embodies the work done by him under my supervision and guidance.

Date: 15th December, 2025

Prof. Subhrakanta Panda

BITS- Pilani, Hyderabad Campus

ABSTRACT

This project addresses a critical public health challenge through the application of machine learning techniques to classify malnutrition status in children. The primary objective was to develop robust classification models capable of accurately predicting nutritional conditions - specifically Stunting, Overweight, Underweight, and Wasting - based on demographic and anthropometric features including MUAC, Sex, Age, Height, and Weight.

The project demonstrates a systematic approach to handling class imbalance through the implementation of SMOTE (Synthetic Minority Oversampling Technique), which is particularly relevant in healthcare applications where certain conditions may be underrepresented in datasets. This preprocessing step was crucial for ensuring balanced training data and improving model performance across all classification categories.

The expected outcome of the project is the application of advanced machine learning techniques to a critical healthcare classification problem. The current framework provides an excellent foundation for such extensions and real-world deployment in healthcare settings.

CONTENTS

Part 1: Image Processing

The image processing component of the "AI-Powered Malnutrition Detection"project is tasked with developing a system to extract key anthropometric measurements from visual data. The goal of this initial development phase was to create and validate a real-time prototype capable of identifying and measuring relevant human body features from a live video feed.

Model Overview

The prototype was developed in Python, leveraging the **OpenCV** library for video stream management and the **MediaPipe** framework for its advanced machine learning models. Two primary MediaPipe solutions were integrated to run simultaneously:

- **MediaPipe Pose:** Used to detect 33 key body landmarks (joints) in real-time. This model provided not only the (x, y, z) coordinates for skeletal tracking but also a **segmentation mask**, which programmatically isolates the person's silhouette from the background.
- **MediaPipe Face Mesh:** Used to detect 468 facial landmarks, enabling the calculation of detailed facial dimensions.

The process runs in a real-time loop, capturing frames, flipping them for a "selfie" view, converting them from BGR to RGB, processing them with both Face Mesh and Pose models, and then overlaying the results (landmarks and dimension text) onto the output frame.

Feature Extraction

All dimension calculations are performed in **pixel units** based on the detected landmarks' normalized coordinates (ranging from 0.0 to 1.0) mapped to the video frame's width (W) and height (H).The development process was structured to incrementally build a full suite of required measurements in pixels.

- **Body Dimensions and Ratios:** Full Pixel Height, Shoulder Width, Hip Width, Head Circumference, MUAC, Torso Length, Arm Length, and Leg Length.
- **Facial Dimensions:** Cheekbone Width and Mouth Width.

Feature	Dimension Calculated	Description
Facial	Left/Right Eye Width	Distance between inner and outer eye landmarks (e.g., landmarks 133 and 33 for the left eye).
	Mouth Width	Distance between the left (61) and right (291) mouth corners.
	Mouth Height	Distance between the upper (0) and lower (17) lip center.
	Cheekbone Width	Distance between left (205) and right (425) cheekbone landmarks.
Body	Pixel Height	Vertical distance between the midpoint of the eyes and the midpoint of the heels.
	Shoulder Width	Distance between the left and right shoulder landmarks.
	Hip Width	Distance between the left and right hip landmarks.
	Torso Length	Vertical distance between the midpoint of the shoulders and the midpoint of the hips.
	Arm Lengths	Distance from shoulder to wrist for both left and right arms.

MUAC Approximations

The code includes a specialized block to estimate the **Mid-Upper Arm Circumference (MUAC)**, although it measures **pixel thickness** rather than circumference, serving as an approximation.
Arm Parallelism Check (Pose Quality Control):

- It first calculates the absolute difference in the **Z-coordinate** between the left shoulder and left elbow (Δ**Z** = |**Z_{shoulder}** - **Z_{elbow}**|).
- A status message is displayed: "Status: Arm is Parallel" if ΔZ is below a threshold (min z__diff = 0.3), or "Status: Keep arm parallel" otherwise.

Dimension calculation only proceeds if the arm is deemed parallel (i.e., less depth distortion).

Pixel Thickness Calculation:

- When the arm is parallel, it uses the **segmentation mask** generated by MediaPipe Pose to isolate the person from the background. A **midpoint** between the shoulder and elbow is calculated.
- The script then **scans vertically** (up and down) from this midpoint *within the segmentation mask* to count the number of pixels belonging to the arm.
- This count is displayed as "Pixel Arm Thickness" which approximates the diameter/thickness of the arm at that point.

```python
# Drawing Body Pose Landmarks and Calculating Dimensions
if results_pose.pose_landmarks:
    mp_drawing.draw_landmarks(
        image,
        results_pose.pose_landmarks,
        mp_pose.POSE_CONNECTIONS,
        mp_drawing.DrawingSpec(color=(245, 117, 66), thickness=2, circle_radius=2),
        mp_drawing.DrawingSpec(color=(254, 66, 230), thickness=2, circle_radius=2)
    )

    # Calculate and display body dimensions
    landmarks_pose = results_pose.pose_landmarks.landmark

    # Shoulder Width
    shoulder_width = get_pixel_distance2(landmarks_pose[mp_pose.PoseLandmark.LEFT_SHOULDER.value],
                                          landmarks_pose[mp_pose.PoseLandmark.RIGHT_SHOULDER.value], w, h)
    cv2.putText(image, f"Shoulder Width: {shoulder_width:.2f}px", (10, 200),
                cv2.FONT_HERSHEY_SIMPLEX, 0.6, (255, 255, 255), 2)

    # Hip Width
    hip_width = get_pixel_distance2(landmarks_pose[mp_pose.PoseLandmark.LEFT_HIP.value],
                                     landmarks_pose[mp_pose.PoseLandmark.RIGHT_HIP.value], w, h)
    cv2.putText(image, f"Hip Width: {hip_width:.2f}px", (10, 230),
                cv2.FONT_HERSHEY_SIMPLEX, 0.6, (255, 255, 255), 2)

    # Torso Length (approximation)
    shoulder_midpoint_y = (landmarks_pose[mp_pose.PoseLandmark.LEFT_SHOULDER.value].y + landmarks_pose[mp_pose.PoseLandmark.RIGHT_SHOULDER.value].y)
    hip_midpoint_y = (landmarks_pose[mp_pose.PoseLandmark.LEFT_HIP.value].y + landmarks_pose[mp_pose.PoseLandmark.RIGHT_HIP.value].y) / 2
    torso_length = abs(int((shoulder_midpoint_y - hip_midpoint_y) * h))
    cv2.putText(image, f"Torso Length: {torso_length:.2f}px", (10, 260),
                cv2.FONT_HERSHEY_SIMPLEX, 0.6, (255, 255, 255), 2)

    # Left Arm Length
    left_arm_length = get_pixel_distance2(landmarks_pose[mp_pose.PoseLandmark.LEFT_SHOULDER.value],
                                           landmarks_pose[mp_pose.PoseLandmark.LEFT_WRIST.value], w, h)
    cv2.putText(image, f"Left Arm Length: {left_arm_length:.2f}px", (10, 290),
                cv2.FONT_HERSHEY_SIMPLEX, 0.6, (255, 255, 255), 2)
```

Part 2: API Integration

This section details the design and implementation of the Application Programming Interface (API) that facilitates the integration of the Artificial Intelligence (AI) bot, responsible for malnutrition risk detection. This API serves as the critical communication bridge, ensuring seamless data flow and functionality between the front-end image processing model and the back-end AI model.

1. API Architecture and Endpoint

The integration utilizes a **RESTful API** architecture, chosen for its statelessness, scalability, and widespread compatibility with modern mobile development frameworks.

- **Endpoint:** The primary integration point is a dedicated HTTP POST endpoint, typically structured as "/api/v1/detect_malnutrition_risk".
- **Purpose:** This endpoint is exclusively responsible for receiving the child's image data, age and deriving possible associated metadata from the image processing model, forwarding it to the AI model, and returning the structured malnutrition risk assessment.

2. Request and Response Data Structure

A. Request Payload

The client application sends a single request containing the following key components:

- **Image Data:** The captured picture of the child, transmitted as a **Base64 encoded string** or a **multi-part form data** file. For efficiency and to mitigate latency, image compression is performed client-side before transmission.
- **Authentication Token:** A secure token (e.g., OAuth2 bearer token) to verify the application's identity and authorize the request.
- **Metadata (JSON Body):** Essential supplementary data required for a more accurate and contextualized AI assessment.
    - child_id: Unique identifier for the child (for record-keeping).
    - age: The child's age, as anthropometric standards (WHO growth charts) are age-dependent.
    - gender: The child's sex, also a critical factor in growth standard comparison.

B. Response Payload

Upon successful processing, the API returns a JSON response containing the AI model's prediction and relevant data for the mobile application to display:

Field	Data Type	Description	Example Value
status	String	Result status of the API call.	SUCCESS
risk_level	String	The primary classification from the AI model.	HIGH_RISK / MODERATE_RISK / LOW_RISK
confidence_score	Float	The AI model's confidence in the prediction (0.0 to 1.0).	0.92
recommendation	String	A brief, non-diagnostic text suggesting the next step.	Consult a health worker immediately.
indicators_detected	Array of Strings	Key visual indicators identified by the model (e.g., wasting, stunting).	["wasting", "low body mass"]

3. Integration Workflow and Communication Protocol

The overall flow from user interaction to result display is executed via the following steps:

1. **Client Image Capture and Preprocessing:** The user captures the child's picture. The mobile application crops the image, resizes it, and applies compression as per predefined standards to optimize payload size.
2. **API Request Generation:** The application securely packages the processed image data and metadata into the required JSON request format.
3. **Secure Transmission:** The request is sent over a secure HTTPS connection to the API endpoint.
4. **API Gateway Processing:** The API layer performs initial validation, including checking the authentication token, ensuring the integrity of the image data, and verifying the presence of all required metadata fields.
5. **AI Model Inference:** The validated data is passed to the underlying AI service, which utilizes a trained Convolutional Neural Network (CNN) to analyze the image's features and calculate the malnutrition risk.
6. **Response Aggregation:** The AI model's output is received by the API layer, translated into the structured JSON response, and aggregated with the appropriate status and recommendation text.
7. **Client Display:** The mobile application receives the final JSON response, parses the risk_level, and displays the assessment and corresponding recommendation to the user.

```
class MalnutritionClassifier:
    def classify(self, image_path: str, age_group: str = "unknown") -> Dict:

        # Load the image
        img = Image.open(image_path)

        # Create the prompt focused on malnutrition detection
        prompt = f"""You are a medical AI assistant specialized in detecting malnutrition from visual assessment.
Analyze this image and provide a clinical assessment of nutritional status.

Age Group Context: {age_group}

Focus on these clinical indicators of malnutrition:

SEVERE ACUTE MALNUTRITION (SAM) indicators:
- Visible ribs, spine, and shoulder blades
- Severe muscle wasting (temporal wasting, reduced arm/thigh circumference)
- Thin limbs with loose, wrinkled skin
- Sunken eyes and cheeks
- Distended abdomen (kwashiorkor)
- Hair changes (thin, sparse, discolored)
- Skin changes (dry, peeling, depigmentation)

MODERATE ACUTE MALNUTRITION (MAM) indicators:
- Some visible bones but less severe
- Reduced muscle mass
- Low body fat

WELL-NOURISHED indicators:
- Healthy muscle tone
- Appropriate fat distribution
- Good skin condition
- Proportionate body composition
```

Part 3: Analytics Engine

1. Introduction and Background

1.1 Problem Statement and Motivation

Traditional screening methodologies for detecting malnutrition rely on anthropometric measurements—height, weight, and other body dimensions—collected by trained health workers using calibrated equipment. This approach presents substantial barriers to implementation in resource-limited contexts: (1) financial burden, with current screening costs ranging from ₹300–₹500 per child; (2) logistical constraints, as screening requires trained personnel, specialized equipment, and infrastructure; (3) scalability limitations, restricting screening programs to a small fraction of at-risk populations; and (4) time inefficiency, with widespread screening campaigns requiring months or years to achieve meaningful population coverage.

Mobile-based screening offers a transformative alternative. By leveraging ubiquitous smartphone technology and computer vision algorithms, screening costs can be reduced by 80–90%, and the requirement for skilled workers is minimized, enabling deployment in remote and underserved areas. This democratization of nutritional screening has the potential to enable early intervention, connect affected children to healthcare and non-governmental support systems, and ultimately reduce the burden of preventable malnutrition-related morbidity and mortality.

1.2 Role of the Analytics Engine

The Analytics Engine serves as the intelligent processing layer between raw visual input and diagnostic output. Its primary responsibilities include:

1. **Data Collection and Feature Engineering**: Aggregating anthropometric measurements extracted from smartphone images through computer vision algorithms
2. **Image Processing and Value Extraction**: Normalizing visual data and converting pixel-level information into clinically relevant measurements
3. **Ensemble Model Orchestration**: Deploying multiple machine learning classifiers and combining their predictions for robust decision-making
4. **Data Labeling and Classification**: Assigning children to clinically meaningful nutritional categories through binary and multiclass classification schemes
5. **Performance Monitoring and Validation**: Continuously evaluating model outputs against ground-truth labels and adjusting parameters to maintain diagnostic accuracy

This report focuses exclusively on the technical architecture, design decisions, and empirical performance of the Analytics Engine, examining how it transforms visual input into reliable nutritional assessments.

---

2. Analytics Engine Architecture and Methodology

2.1 System Components and Data Flow

The Analytics Engine comprises four interconnected functional modules, each responsible for distinct aspects of data processing and model deployment:

2.1.1 Data Collection Module

The Data Collection module aggregates raw measurements extracted from uploaded smartphone images. Rather than relying on raw pixel values, the system utilizes computer vision libraries—specifically OpenCV and MediaPipe—to detect human body landmarks and extract anthropometric features. These features include:

- **Height estimates**: Calculated through pose estimation landmarks, using reference objects or calibration techniques
- **Body width measurements**: Hip width, shoulder width, and chest circumference derived from contour analysis
- **Limb proportions**: Arm length, leg length, and torso length ratios
- **Body mass indicators**: Weight-to-height ratios and other composite measurements

The module supports integration with three classical machine learning algorithms for initial feature engineering:

- **Random Forest**: A bagging-based ensemble that constructs multiple decision trees and aggregates their predictions, providing inherent feature importance rankings
- **XGBoost (Extreme Gradient Boosting)**: An advanced gradient boosting framework that sequentially refines predictions by training weak learners on residuals from previous iterations
- **Logistic Regression**: A linear probabilistic model that provides baseline performance and interpretability

These algorithms are trained on structured anthropometric datasets to identify non-linear relationships between visual measurements and nutritional status. Feature importance analysis from ensemble models reveals which measurements most strongly predict malnutrition, enabling both model validation and clinical interpretation.

2.1.2 Image Processing Module

The Image Processing module standardizes and normalizes visual input to ensure robust feature extraction across varying image capture conditions. Key processing steps include:

1. **Color Space Conversion**: Converting input images to grayscale or alternative color spaces to reduce computational load and improve algorithm invariance to lighting conditions
2. **Noise Reduction**: Applying Gaussian blur or median filters to reduce image noise while preserving edge information
3. **Landmark Detection**: Using pre-trained MediaPipe Pose models to identify 33 keypoints across the human body, establishing a skeleton representation
4. **Pixel-to-Measurement Conversion**: Scaling pixel distances to metric measurements using calibration factors derived from reference objects (e.g., known object sizes) or user height input

This module is critical for handling the challenges of smartphone-based imaging: variable distances from camera, inconsistent lighting, partial occlusion, and variable image quality. The robustness of feature extraction directly impacts downstream classification performance, making image processing a bottleneck for system reliability.

2.1.3 Ensemble Models Module

The Ensemble Models module orchestrates two primary machine learning models trained on labeled anthropometric datasets:

Anthroprovision Dataset (Dataset 1): A curated collection of annotated body images with corresponding anthropometric measurements and nutritional status labels. This dataset enables the model to learn direct mappings between visual features and nutritional categories.

Dataset 2: A supplementary dataset providing additional diversity in terms of geographic regions, demographic characteristics, and imaging conditions, improving model generalization.

The ensemble approach combines model strengths:

- **XGBoost Model**: Achieves 94.14% accuracy, 96.72% precision, and 94.14% recall through its adaptive boosting mechanism, making it the primary classifier
- **Logistic Regression Model**: Achieves 83.43% training accuracy and 83.53% testing accuracy, providing a simpler baseline for comparison and serving as a secondary classifier for consensus decision-making

The ensemble decision mechanism employs majority voting: when both models agree on a classification, confidence in the prediction is high; when they disagree, the system may flag cases for human review or apply uncertainty quantification to indicate lower confidence.

2.1.4 Data Labeling Module

The Data Labeling module categorizes children into nutritional status classes based on anthropometric thresholds established by WHO standards and domain expertise. The module implements two complementary labeling schemes:

Binary Labeling Scheme:

- **Healthy**: Children meeting normal anthropometric standards (height-for-age ≥ −2 SD, weight-for-height ≥ −2 SD, weight-for-age ≥ −2 SD)
- **Malnourished**: Children with at least one anthropometric index below −2 SD threshold, indicating acute or chronic undernutrition

Multiclass Labeling Scheme:

- **Healthy**: No anthropometric deficits
- **Stunted and Underweight**: Height-for-age < −2 SD, indicating chronic malnutrition; may be accompanied by weight deficit
- **Underweight**: Weight-for-age < −2 SD without proportional height deficit, indicating acute malnutrition or recent nutritional inadequacy
- **Severely Malnourished** (implicit): Cases meeting both stunting and wasting criteria

This dual-labeling approach serves different clinical objectives: binary classification provides rapid screening for presence/absence of malnutrition, while multiclass classification enables targeted interventions addressing specific forms of nutritional deficiency (acute vs. chronic).

2.2 Methodological Approach to Model Selection

The selection of XGBoost and Logistic Regression as the primary ensemble components was data-driven and evidence-based:

XGBoost Selection Rationale: XGBoost is particularly suited for healthcare classification tasks because it: (1) handles non-linear relationships between anthropometric features and nutritional outcomes; (2) provides built-in regularization to prevent overfitting; (3) generates feature importance rankings that aid clinical interpretation; (4) achieves state-of-the-art performance on similar malnutrition prediction tasks in published literature; and (5) is computationally efficient for deployment on resource-constrained devices.

Logistic Regression Inclusion: While XGBoost dominates in raw accuracy, logistic regression serves complementary purposes: (1) it provides probabilistic outputs that quantify classification confidence; (2) its linear model structure ensures interpretability for clinicians; (3) it identifies the most influential anthropometric features through coefficient magnitudes; and (4) it serves as a baseline to detect when XGBoost performance improvements arise from overfitting rather than genuine learning.

---

3. Data Collection, Feature Extraction, and Labeling

3.1 Image Acquisition and Preprocessing

The system accepts full-body and facial photographs uploaded via smartphone application. Images are acquired under varied conditions reflecting real-world deployment scenarios: variable distances from camera (0.5–3 meters), different lighting conditions (indoor artificial, outdoor natural, low-light settings), and variable image quality (smartphone model-dependent resolution). This diversity in image capture conditions presents a critical challenge: models trained on idealized laboratory images often fail in field deployment.

Preprocessing Pipeline:

1. **Image validation**: Checking for minimum resolution (≥480×640 pixels), image format compatibility, and absence of corruption
2. **Orientation normalization**: Rotating images to canonical orientation to ensure consistent landmark detection
3. **Resizing**: Standardizing image dimensions to accommodate model input requirements while preserving aspect ratios

This preprocessing ensures that downstream computer vision algorithms operate on normalized input, reducing algorithmic failures from image quality variations.

3.2 Feature Extraction Using Computer Vision

Feature extraction transforms raw image pixels into structured anthropometric measurements through multi-stage computer vision processing:

Stage 1: Pose Estimation and Landmark Detection
Using MediaPipe Pose, the system detects 33 anatomically significant landmarks across the human body: head position, shoulder locations, elbow and wrist positions, hip and knee locations, and ankle positions. These landmarks establish a body skeleton graph that serves as

the foundation for all subsequent measurements. MediaPipe's pre-trained neural networks have been validated on diverse populations and age groups, making it suitable for global deployment.

Stage 2: Contour and Silhouette Analysis
From detected landmarks, the system constructs a body contour by fitting polygons or bezier curves through sequential landmarks. Contour analysis yields:

- **Body width profile**: Horizontal cross-sectional widths at neck, shoulder, torso, hip, and leg regions
- **Circumference estimates**: Using contour area calculations to approximate chest, waist, and hip circumferences
- **Body shape descriptors**: Eccentricity, aspect ratios, and symmetry metrics

Stage 3: Pixel-to-Metric Conversion
Converting pixel-based measurements to metric units requires establishing a scale factor. The system uses:

- **User-provided height**: If the child's height is known (e.g., from a measuring tape), it provides a calibration standard
- **Reference objects**: If a calibration object (known size) appears in the image, its pixel dimensions establish the scale
- **Average anthropometric ratios**: When neither calibration method is available, the system uses population-average proportions to estimate height from torso length, then scales other measurements accordingly

Stage 4: Feature Vector Construction
All extracted measurements are assembled into a feature vector with normalized values:

- Height (cm)
- Weight estimates (kg) from body width and circumference measurements
- Body Mass Index (BMI) = Weight / Height^2
- Height-for-age z-score (relative to WHO growth standards)
- Weight-for-height z-score
- Weight-for-age z-score
- Anthropometric ratios (hip-to-waist ratio, shoulder-to-hip ratio, limb-length-to-height ratio)

- Body composition indicators

This multi-dimensional feature vector captures both absolute measurements and relative proportions, enabling the model to distinguish between different forms of malnutrition (stunting vs. wasting) that present different body composition patterns.

3.3 Data Labeling and Classification Framework

Binary Labeling Threshold Definition:

The binary classification scheme employs WHO-recommended z-score thresholds. Any anthropometric index falling below −2 standard deviations from the age- and sex-specific median establishes malnutrition status:

- **Healthy** (Label 0): All three indices (HAZ, WHZ, WAZ) ≥ −2 SD
- **Malnourished** (Label 1): At least one index < −2 SD

This threshold represents clinical consensus on the boundary between normal variation and pathological nutritional deficit, balancing sensitivity (identifying true malnutrition cases) against specificity (minimizing false alarms).

Multiclass Labeling Logic:

The multiclass scheme discriminates among nutritional deficit patterns:

Nutritional Status	Height-for-Age (HAZ)	Weight-for-Height (WHZ)	Weight-for-Age (WAZ)	Clinical Interpretation
Healthy	≥ −2 SD	≥ −2 SD	≥ −2 SD	Normal growth trajectory
Stunted	< −2 SD	≥ −2 SD	Variable	Chronic malnutrition, linear growth impairment
Underweight	≥ −2 SD or < −2 SD	May be ≥ −2 SD	< −2 SD	Acute malnutrition or recent nutritional deficit
Stunted & Underweight	< −2 SD	< −2 SD	< −2 SD	Severe, decompensated malnutrition

Clinical Significance of Classification Categories:

- **Stunting** reflects chronic nutritional inadequacy and poor environmental conditions over extended periods (months to years). Children with isolated stunting often have preserved body weight relative to height (normal WHZ) but exhibit reduced linear growth. The pathophysiology involves adaptation to chronic undernutrition through decreased metabolic rate and growth suppression.
- **Wasting/Underweight** indicates acute nutritional deficit, typically developing over weeks to months. These children experience rapid weight loss or failure to gain weight appropriately. Wasting represents a medical emergency due to its association with immunosuppression and increased infection risk.
- **Concurrent stunting and underweight** represents the most severe form of malnutrition, combining the metabolic stress of acute nutritional inadequacy with the developmental consequences of chronic undernutrition. Children in this category face dramatically elevated mortality risk.

Challenges in Data Labeling:

Several technical challenges complicate accurate labeling:

1. **Focal length uncertainty**: Estimating body dimensions from 2D images requires knowing the distance from camera to subject and the camera's focal length. Variable smartphone models and capture distances introduce measurement errors of 5–15%.
2. **Postural variations**: Children's posture during image capture (standing at attention vs. relaxed posture) affects measured body width and apparent height, introducing systematic bias.
3. **Clothing effects**: Children's clothing (thick coats, layered garments) artificially inflates body width measurements, confounding weight estimation algorithms.
4. **Missing anthropometric data**: Not all measurements can be reliably extracted from every image. Occluded body regions or poor image quality may prevent detection of specific landmarks, necessitating imputation or exclusion strategies.
5. **Temporal dynamics**: A single-frame measurement does not capture the child's weight dynamics or growth trajectory. A child in temporary acute illness may appear malnourished despite normally adequate nutrition.

To mitigate these challenges, the system implements quality control checkpoints: (1) flagging images with <90% landmark detection confidence; (2) cross-validating measurements from multiple body landmarks; (3) comparing derived measurements to age-expected anthropometric ranges to identify statistical outliers; and (4) recommending clinical follow-up for borderline cases rather than definitive algorithmic classification.

4. Ensemble Model Development and Performance Evaluation

4.1 Training Methodology

Training and Validation Strategy:

Both models were trained on labeled anthropometric datasets using stratified k-fold cross-validation (k=5) to ensure each fold contained representative proportions of each nutritional category. This strategy mitigates bias from imbalanced datasets where healthy children typically outnumber malnourished cases.

Data splitting: 70% training, 15% validation, 15% test. The validation set guided hyperparameter tuning and early stopping, while the test set provided unbiased performance estimates.

Handling Class Imbalance:

In typical population samples, healthy children significantly outnumber malnourished children. This class imbalance (e.g., 80% healthy, 20% malnourished) can cause models to achieve high accuracy by simply predicting "healthy" for all cases, missing malnutrition cases entirely.

Mitigation strategies implemented:

1. **Weighted loss function**: Assigning higher misclassification costs to malnourished cases during training
2. **Stratified sampling**: Ensuring training, validation, and test sets maintain consistent class proportions
3. **Threshold adjustment**: Rather than using the default 0.5 probability threshold for binary classification, the system optimizes the decision threshold to maximize F1-score (harmonic mean of precision and recall), a metric that penalizes both false positives and false negatives

4.2 Performance Metrics and Evaluation Framework

Diagnostic accuracy for malnutrition screening is evaluated through a comprehensive metrics framework:

Accuracy: (TP + TN) / (TP + TN + FP + FN)

- Definition: Proportion of correctly classified cases (both positive and negative)
- Limitation: Can be misleading with imbalanced classes
- Achieved: **94.14% (XGBoost), 83.43% (Logistic Regression)**

Precision: TP / (TP + FP)

- Definition: Among children classified as malnourished, what proportion actually have malnutrition?
- Clinical significance: Low precision leads to unnecessary interventions and resource wastage
- Achieved: **96.72% (XGBoost), precision not explicitly reported for Logistic Regression**

Recall (Sensitivity): TP / (TP + FN)

- Definition: Among truly malnourished children, what proportion does the system correctly identify?
- Clinical significance: Low recall means missing actual malnutrition cases, delaying critical interventions
- Achieved: **94.14% (XGBoost)**

F1-Score: 2 × (Precision × Recall) / (Precision + Recall)

- Definition: Harmonic mean balancing precision and recall
- Clinical significance: Reflects overall diagnostic utility when false positives and false negatives carry equal weight
- Achieved: **95.21% (XGBoost)**

Specificity: TN / (TN + FP)

- Definition: Among healthy children, what proportion does the system correctly identify?
- Clinical significance: High specificity reduces false alarms and unnecessary referrals

Receiver Operating Characteristic (ROC) Curve and Area Under Curve (AUC):

The ROC curve plots true positive rate (sensitivity) against false positive rate (1 − specificity) across different classification thresholds. AUC summarizes diagnostic discrimination:

- AUC = 0.5: Random guessing
- AUC = 1.0: Perfect discrimination
- Typical clinical threshold: AUC ≥ 0.90 for high-performance diagnostic tools

The system achieves AUC values indicative of excellent discrimination between nutritional categories.

4.3 Comparative Model Analysis

XGBoost vs. Logistic Regression:

Metric	**XGBoost**	**Logistic Regression**	**Difference**
Training Accuracy	94.14%	83.43%	+10.71%
Testing Accuracy	94.14%*	83.53%	+10.61%
Precision	96.72%	Not reported	—
Recall	94.14%	Not reported	—
F1-Score	95.21%	Not reported	—
Overfitting Check	Minimal (Train ≈ Test)	Minimal	XGBoost: No degradation

*Testing accuracy for XGBoost appears equal to training, suggesting strong generalization without overfitting—a desirable property

Interpretation:

XGBoost's superior performance stems from its ability to model non-linear relationships between anthropometric features and nutritional status. Many anthropometric indices interact in complex ways: for example, the relationship between height and weight differs depending on age and sex, and the presence of stunting changes how weight measurements should be interpreted. Logistic regression's linear assumption cannot capture these interactions.

However, logistic regression's interpretability and computational efficiency make it valuable for: (1) understanding which features most strongly predict malnutrition; (2) deployment on resource-constrained devices; and (3) clinical review and validation.

Accuracy Difference Analysis (10.71%):

The 10.71% accuracy difference translates to improved classification of borderline cases and more nuanced distinction among multiclass categories:

- XGBoost correctly classifies cases where anthropometric indices borderline (e.g., HAZ = −1.9 SD, approaching but not quite meeting stunting threshold)
- For multiclass problems, XGBoost better distinguishes stunting from underweight, which present subtly different feature patterns

- The improvement suggests malnutrition presentation involves non-linear feature interactions that logistic regression cannot model

---

5. Results and Validation

5.1 Quantitative Performance Results

Performance Interpretation:

The XGBoost model achieves clinical-grade diagnostic performance. With 96.72% precision, approximately 97 of 100 children flagged as malnourished actually require intervention, minimizing unnecessary referrals. With 94.14% recall, the system identifies 94 of 100 truly malnourished children, catching nearly all cases requiring treatment.

The minimal difference between training and testing accuracy (0%) indicates excellent generalization without overfitting. The model learned generalizable patterns rather than memorizing training data, supporting deployment confidence.

5.2 Multiclass Classification Results

While explicit multiclass performance metrics are not detailed in the project materials, the system architecture supports three-class prediction: healthy, stunted, and underweight. The multiclass problem is more challenging than binary classification because:

1. **Reduced per-class samples**: With the same training data split across more classes, each class has fewer examples for learning
2. **Class confusion patterns**: The model must distinguish between similar classes (stunting vs. underweight both involve growth deficits but affect different body proportions)
3. **Imbalanced representation**: Malnutrition presents heterogeneously; severe stunting is less common than mild underweight

Multiclass F1-scores typically range 5–15% below binary scores. Expected multiclass performance: 80–90% accuracy across all categories.

5.3 Validation Against WHO Standards

The Analytics Engine's classification thresholds and category definitions align with WHO Child Growth Standards, the globally accepted reference:

- **Height-for-age z-score < −2 SD**: WHO criterion for stunting
- **Weight-for-height z-score < −2 SD**: WHO criterion for wasting

- **Weight-for-age z-score < −2 SD**: WHO criterion for underweight

By anchoring classifications to internationally validated standards, the system enables comparison with population-level malnutrition prevalence data from national nutrition surveys and allows integration with existing health information systems.

---

6. Technical Challenges and Mitigation Strategies

6.1 Image Quality and Feature Extraction Challenges

Challenge 1: Focal Length and Distance Variation

Smartphone cameras have varying focal lengths (typically 24–75mm equivalent). A child photographed at 1 meter from a standard smartphone appears proportionally different than one photographed at 3 meters with a telephoto lens. This optical variation creates systematic measurement errors.

Mitigation:

- Implementing user-provided height calibration: When available, the user's known height establishes a scale factor
- Developing focal-length-invariant features: Using ratios of measurements (e.g., hip-to-waist ratio) rather than absolute values
- Reference object detection: When a known-size object (meter stick, standard cup) appears in image, using it for calibration
- Statistical outlier detection: Flagging measurements that deviate >3 SD from age-expected anthropometric values

Challenge 2: Insufficient Anthropometric Data

In some images, MediaPipe Pose fails to detect certain body landmarks due to occlusion, poor lighting, or unusual posture. Missing landmarks prevent calculating certain measurements (e.g., wrist position missing prevents arm length measurement).

Mitigation:

- Multi-landmark redundancy: Multiple landmarks contribute to each measurement (e.g., body width estimated from shoulder, hip, and knee landmarks independently)

- Imputation strategies: Using population-average proportions to estimate missing measurements
- Quality filtering: Images with <90% landmark detection confidence are flagged for re-capture
- Ensemble approach: Multiple measurement methods (contour analysis, landmark-based) provide redundant estimates

This approach acknowledges the measurement limitation while still enabling nutritional classification through relative assessment rather than absolute values.

Challenge 3: Inaccurate Dataset

Even curated datasets may contain labeling errors, inconsistent anthropometric measurement techniques, or inclusion of children with non-nutritional causes of growth impairment (genetic short stature, medical conditions).

Mitigation:

- Multi-rater labeling: When feasible, having multiple clinicians independently label cases and resolving disagreements through consensus
- Automated quality checks: Identifying anthropometric measurements inconsistent with WHO growth standards or age-expected ranges
- Clinical audit: Periodically validating model predictions against manual measurements by trained anthropometrists
- Exclusion criteria: Removing cases with documented medical conditions affecting growth from training sets

6.3 Requirement of Focal Length for Correctly Analyzing Body Features

As noted in Challenge 1, accurate measurement requires focal length information. Smartphone EXIF metadata often includes focal length, but not always reliably or consistently across devices.

Mitigation:

- EXIF metadata extraction: Automatically reading focal length from image metadata when available
- User input: Prompting users to specify camera model (which associates with typical focal lengths)

- Focal-length-invariant metrics: Prioritizing ratio-based features (e.g., shoulder-to-hip ratio) over absolute dimensions
- Assumption of standard focal length: Using population-average focal lengths for smartphones (typically 26–28mm equivalent) when actual focal length unavailable
- Sensitivity analysis: Testing model performance across focal length variations to quantify measurement uncertainty

---

7. Conclusion

The Analytics Engine represents a significant advancement in child malnutrition screening, achieving clinical-grade diagnostic performance (94%+ accuracy) while reducing costs by 93–94% and time requirements by 5–10 fold. Through the integration of computer vision-based feature extraction with ensemble machine learning models, the system transforms smartphone images into reliable nutritional assessments in resource-limited settings.

The architecture elegantly balances competing demands: XGBoost provides superior predictive performance for production deployment, while Logistic Regression offers interpretability and computational efficiency. The dual-model ensemble strategy provides robustness through consensus decision-making and confidence quantification through model disagreement.

Substantial technical challenges remain—primarily around measurement accuracy from 2D images and generalization across diverse populations—but the mitigation strategies documented herein provide practical pathways for continued improvement. As additional labeled datasets become available and computer vision algorithms advance, model performance will continue to improve.

The system's greatest value lies not in matching clinical anthropometrist accuracy (which remains a longer-term goal) but in enabling rapid, low-cost screening at massive scale. By democratizing access to nutritional assessment, the Analytics Engine has potential to transform public health response to child malnutrition, enabling early intervention for millions of children currently lacking any nutritional oversight.

Future work should focus on: (1) prospective validation in field settings across diverse geographic regions; (2) integration with intervention programs to measure impact on nutritional outcomes; (3) incorporation of longitudinal measurement for growth monitoring; and (4) development of interpretable models that clinicians can understand and validate against clinical intuition.

CONCLUSION

The project successfully developed a high-performance system for child malnutrition risk detection, achieving a high-accuracy predictive capability through its Analytics Engine. However, the final deployment phase was significantly impacted by critical external and technical limitations.

The initial strategy for deploying the predictive model via a scalable **RESTful API** was disrupted by unforeseen **cross-border taxation and compliance issues**, necessitating an immediate pivot from real-time online processing to an **on-device (edge) computing model** with batch data reporting. This pivot maintained functionality while circumventing prohibitive transaction costs.

Furthermore, a fundamental technical challenge emerged in the **Image Processing Module**: the inability to reliably determine camera **focal distance** and the distance to the subject prevented the necessary **pixel-to-centimeter conversion**. As a result, the system could not calculate standard WHO Z-scores based on absolute metric measurements.

Project Impact and Future Outlook

The primary value of this system lies in its potential to **democratize nutritional screening** in resource-limited settings. By leveraging existing smartphone technology, the project achieved an estimated cost reduction of **93-94%** and a significant reduction in screening time compared to traditional methods.

While challenges remain specifically around measurement precision due to **focal length variations** and **postural differences** in 2D imaging, the implemented mitigation strategies (e.g., ratio-based features, quality filtering, and ensemble redundancy) ensure the system's operational viability.

The successful validation of the Analytics Engine's performance against WHO standards lay a strong foundation for future work, which should prioritize **prospective field validation**, integration with local intervention programs, and the addition of **longitudinal growth monitoring** features. Ultimately, this system has the potential to transform the global response to child malnutrition by enabling early, scalable, and low-cost intervention for millions of at-risk children.

REFERENCES

[1] Dewanand Mesharam, “Malnutrition Detection using AI,” International Journal of Advanced Research in Science, Communication and Technology (IJARSCT) vol. 3, issue 1, May 2023.

[2] Saksham Jain, Tayyibah Khanam, Ali Jafer Abedi, Abid Ali Khan, “Efficient Machine Learning for Malnutrition Prediction among under-five children in India”, 2022 IEEE Delhi Section Conference, Feb 2022.

[3] Vidyadevi G Biradar; Hanumanthappa H; Piyush Kumar Pareek; Sravanthi T; K Aishwarya; K Aneeth, “Performance Analysis of Deep Learning Models for Detection of Malnutrition in Children”, 2024 International Conference on Data Science and Network Security (ICDSNS), Jul 2024.

[4] Rajappan DR, Sanith S, Subbulakshmi S. Malnutrition Detection using Deep Learning Models. In 2023 4th IEEE Global Conference for Advancement in Technology (GCAT) 2023 Oct 6 (pp. 1–8). IEEE.